

GARIL AI

Subprocessors & Data Flows

Third parties that may process institutional data

Version	1.0
Date	3 October 2026
Audience	Partner universities & institutional stakeholders
Classification	Confidential — authorised institutional use

Important notice

Illustrative schedule based on the standard GARIL AI architecture. Exact vendors and regions depend on the deployment (self-hosted university VPS vs managed hosting). Update this schedule in the signed DPA.

1. Data flow (standard deployment)

- User browser!' TLS reverse proxy / CDN!' GARIL AI web & API
- API!' MongoDB (accounts, research, governance data)
- API!' optional S3-compatible object storage (uploads/assets)
- API!' LLM provider (prompts/context for AI features)
- API!' literature/search APIs as configured (e.g. OpenAlex, PubMed, Tavily)
- API!' transactional email provider (password reset / notices)

2. Typical subprocessors

Function	Typical provider	Data involved
Hosting / VPS	Institution or designated host (e.g. Coolify/VPS)	Full application stack & DB
AI inference	OpenRouter (routes to model vendors)	Prompts, context, outputs, usage
Email	Resend (or configured SMTP/ESP)	Name, email, reset links
Object storage	MinIO/S3-compatible (optional)	Uploaded files & metadata
Literature APIs	OpenAlex, NCBI/PubMed, others as enabled	Search queries; generally non-account PII

3. University control

- Feature flags can disable AI or modules per university
- Institutions should instruct users not to paste unnecessary special-category data into prompts
- Self-hosted deployments keep primary data on university-controlled infrastructure

4. Change management

Material subprocessor changes will be communicated to institutional contacts as required by the DPA (typically advance notice for non-emergency changes).